

CS143 – HW3

*PART 1*

```
SELECT company,  
SUM(value='agile-dev') AS AgileDev,  
SUM(value='benefit-company') AS BenefitCompany,  
SUM(value='bonded-by-product') AS BondedByProduct,  
SUM(value='continuous-delivery') AS ContinuousDelivery,  
SUM(value='creative-innovation') AS CreativeInnovation,  
SUM(value='cross-dep') AS CrossDep,  
SUM(value='customer-first') AS CustomerFirst,  
SUM(value='data-driven') AS DataDriven,  
SUM(value='diverse-team') AS DiverseTeam,  
SUM(value='engages-community') AS EngagesCommunity,  
SUM(value='engineering-driven') AS EngineeringDriven,  
SUM(value='eq-iq') AS EqIq,  
SUM(value='fast-paced') AS FastPaced,  
SUM(value='feedback') AS Feedback,  
SUM(value='flat-organization') AS FlatOrganization,  
SUM(value='flex-hours') AS FlexHours,  
SUM(value='friends-outside-work') AS FriendsOutsideWork,  
SUM(value='good-beer') AS GoodBeer,  
SUM(value='impressive-teammates') AS ImpressiveTeammates,  
SUM(value='inclusive') AS Inclusive,  
SUM(value='internal-mobility') AS InternalMobility,  
SUM(value='internal-promotion') AS InternalPromotion,  
SUM(value='interns') AS Interns,  
SUM(value='junior-devs') AS JuniorDevs,  
SUM(value='light-meetings') AS LightMeetings,  
SUM(value='lunch-together') AS LunchTogether,  
SUM(value='many-hats') AS ManyHats,  
SUM(value='new-tech') AS NewTech,  
SUM(value='office-layout') AS OfficeLayout,  
SUM(value='open-communication') AS OpenCommunication,  
SUM(value='open-source') AS OpenSource,  
SUM(value='pair-programs') AS PairPrograms,  
SUM(value='parents') AS Parents,  
SUM(value='personal-growth') AS PersonalGrowth,  
SUM(value='physical-wellness') AS PhysicalWellness,  
SUM(value='product-driven') AS ProductDriven,  
SUM(value='project-ownership') AS ProjectOwnership,  
SUM(value='psychologically-safe') AS PsychologicallySafe,  
SUM(value='quality-code') AS QualityCode,  
SUM(value='rapid-growth') AS RapidGrowth,  
SUM(value='remote-ok') AS RemoteOk,  
SUM(value='retention') AS Retention,
```

```

SUM(value='risk-taking') AS RiskTaking,
SUM(value='safe-env') AS SafeEnv,
SUM(value='team-oriented') AS TeamOriented,
SUM(value='worklife-balance') AS WorklifeBalance
FROM keyvalues
GROUP BY company;

```

Number of rows = 67  
 Number of columns = 47

## PART 2

### 1. NO Subquery with DISTINCT

id	select_type	table	partitions	type	possible_keys	key
	key_len	ref	rows	filtered	Extra	
1	SIMPLE	caltrans	NULL	ALL	PRIMARY	NULL
	NULL NULL	103449	20.99	Using where;	Using temporary;	Using filesort

### 2. SELECT within a SELECT

id	select_type	table	partitions	type	possible_keys	key
	key_len	ref	rows	filtered	Extra	
1	PRIMARY	<derived2>	NULL	ALL	NULL NULL	NULL
	NULL 2411	100.00	Using temporary;	Using filesort		
2	DERIVED	caltrans	NULL	ALL	PRIMARY	NULL
	NULL NULL	103449	2.33	Using where;	Using temporary;	Using filesort

### 3. JOIN as a FILTER

id	select_type	table	partitions	type	possible_keys	key
	key_len	ref	rows	filtered	Extra	
1	PRIMARY	<derived2>	NULL	ALL	NULL NULL	NULL
	NULL 24129	100.00	Using temporary;	Using filesort		
2	DERIVED	c	NULL	ALL	PRIMARY	NULL
	NULL 103449	2.33	Using where;	Using temporary;	Using filesort	
2	DERIVED	<derived3>	NULL	ref	<auto_key0>	
	<auto_key0> 8	hw2.c.highway	10	100.00	NULL	
3	DERIVED	caltrans	NULL	ALL	PRIMARY	NULL
	NULL NULL	103449	2.33	Using where;	Using temporary	

### 4. USING an IN subquery as a Filter

id	select_type	table	partitions	type	possible_keys	key
	key_len	ref	rows	filtered	Extra	

1	PRIMARY	<derived2>	NULL	ALL	NULL	NULL	NULL	NULL
		NULL 2411	100.00	Using temporary;	Using filesort			
2	DERIVED	c	NULL	ALL	PRIMARY	NULL	NULL	NULL
		NULL 103449	2.33	Using where;	Using temporary;	Using filesort		
2	DERIVED	<subquery3>	NULL	eq_ref	<auto_key>			
		<auto_key> 265	hw2.c.highway, hw2.c.area	1	100.00	NULL		
3	MATERIALIZED	caltrans	NULL	ALL	NULL	NULL	NULL	NULL
		NULL NULL 103449	2.33	Using where				

## 5. FORMAL Left Semijoin

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra		
1	PRIMARY	<derived2>	NULL	ALL	NULL	NULL	NULL	NULL	100.00	Using temporary;	Using filesort		
2	DERIVED	c	NULL	ALL	PRIMARY	NULL	NULL	NULL	103449	2.33	Using where;	Using temporary;	Using filesort
3	SUBQUERY	caltrans	NULL	ALL	NULL	NULL	NULL	NULL	103449	2.33	Using where		

From the outputs we notice that query 1 included no sub queries or derived tables. Query 2 has 1 derived table, Query 3 has 3 derived tables, Query 4 has 2 derived tables and Query 5 has 1 sub query and 1 derived table. The column type in the above outputs represents the type of join. In the case of MySQL this is always ALL to represent a Cartesian join. The column rows filtered shows us the percentage of rows it actually filtered. The ref column shows the number of rows to actually filter. The last extra column represents extra information about the query for instance 'using where' means that the query had a where clause, 'filesort' refers to some kind of sort, 'temporary' refers to having a temporary table etc. We notice that Query 3 has a large number of rows it filters and also has temporary tables. These not only take up more memory but also take longer and thus are more inefficient. Similarly, Query 4 filters many rows as well. We see that Query 2 is pretty efficient (more than Query 1) because it uses its derived table to filter rows quicker and is thus more efficient. This is the same for Query 5.

### PART 3

1.  
R is scanned 100,000 times because it goes through the loop each time.
2.  
a) Each tuple in L and R is scanned once.  
b)  
in the worst case, number of block transfers would be

number of blocks in L + (height of the B+ tree of index R+!) \* (number of tuples in L)

3.

We prefer Block Nested Loop Join if the outer table contains fewer blocks in memory.

We prefer the Naïve Nested loop Join in conditions where the block size is almost equal to the tuple size. This is also preferred when no index is used on the join.

We prefer the Indexed Nested Loop Join when R is indexed accurately and the length is small so the time it takes is less.

4.

If a is a tuple in R1 and b is a tuple in R2. But since we have R2-R3, b is also a tuple in R3.

Thus tuple x will show up in R1 join (R2-R3) if a is in R1, b is in R2 and there is a b isn't in R3.

The tuple x again will show up in (R1 join R2) – (R2 join R3) if a is in R1 and b is in R2 and NOT (a is in R1 and b is in R3)

But as we know from above, a is in R1, we can say that the above two conditions are the same.

Thus this equation holds true.

This can help make the joins more efficient because we can reduce the number of fields before doing the join (which is tedious).

5.

Consider two relations R and S as shown below:

R:

A	B
1	2
2	3
3	4

S:

A	B
1	6
2	3

7	8
---	---

For these set of relations,

$$\pi_A(R - S) =$$

A
1
3

$$\text{But, } \pi_A(R) - \pi_A(S) =$$

A
3

In this instance, the equation doesn't hold.